

Stage 11: OSB (Operationally Sufficient Basin) Checklist

Purpose: You don't need a mathematically perfect single basin. You need a basin good enough to guarantee stability and low hallucination — the denoiser stack does the rest.

Hard Behavior Gates (must pass)

Accuracy exact	≥ target (LatentARC: 100/100 exact)
Hallucination	≈ noise floor (≤0.5%)
Token drift stability	r_trend_tokens ≥ 0.90
Trajectory health	Radius ↓, SNR ↑ over steps

Soft Landscape Health (advisory, not blocking synthetic)

Phantom Index (PI)	≤ 0.30 postwarp
Margin	≥ 0.03 (true well deeper than runnerup)
S_median	≥ 0.50 (apex sharpness)

Stress Probes (fast checks)

Pass ≥3/4 while still meeting hard gates:

- Half the steps
- +25% jitter
- Tap shift (−3 ↔ −2)
- 20% prompt mix shift

When to Tighten Bars (LLM integration, not synthetic)

Phantom Index (PI)	≤ 0.15
Margin	≥ 0.04
S_median	≥ 0.55
Calibration	≥ 300 ontopic calibration prompts

Bottom Line

- Clean is nice. Good enough is plenty: if hallucination is suppressed and trajectories stabilize, you've met Stage 11's doctrine (Warp → Detect → Denoise).
- Use PI/margin as dashboard lights, not blocking bars.