Seed Document — Stage 11 A/B/C Experimental Program

---

0) North Star

The doctrine *Warp → Detect → Denoise* is frozen in patents and math (Appendix A) and now baselined in practice (v4b config). The A/B/C experiments are how we move from **latent promise** to **product-grade validation**: do we see measurable, repeatable uplifts in prompt handling when the full doctrine is active inside an LLM?

---

1) What We Have

- **Patents & Math**: Energy-well warping, phantom suppression, calibration, denoising (Appendix A, patent filings)【539†patent_appendix_a.pdf】.
- **Integration Readiness**: Layer-9 wiring plan, hook mechanics, calibration profiles (LLM integration plan)【533†stage11_llm_integration_plan.pdf】.
- **Baseline**: Always-on warp (geo mode) at layer −9, tuned for budget GPUs (T4/L4). v4b config passes pass/fail checks and is now frozen【543†Stage-11 Warp — Wrap-up & Baseline V4b.pdf】.
- **A/B Results**: First wobble-pack tests show neutral-to-positive deltas, stable application rates, and string-level improvements in wobble exits【542†ab_results_geo_base.json】.

---

2) A/B/C Framework

A/B tests: Stock vs Warp. Already crossed threshold → warp has measurable effect.

B/C tests: Warp vs Warp+Detect, Warp+Detect vs Warp+Detect+Denoise. - **B (Warp+Detect)**: Ensures boosts localize where drift evidence exists (soft precision). - **C (Warp+Detect+Denoise)**: Confirms stability and phantom suppression when doctrine runs in full.

Together A/B/C forms the hinge point: we prove incremental value at each layer of the doctrine.

---

3) Why It Matters

- **Patents**: Doctrine is already protected. A/B/C experiments are how we demonstrate *working reduction to practice*.
- **Narrative**: Warp alone improves geometry; Detect makes it precise; Denoise makes it reliable. Each step builds the case for productization.
- **Baseline to Product**: With v4b frozen, deltas from B and C will show whether the doctrine scales into a plug-in runtime module (continuation claim optional).

---

4) What to Watch

- **Metrics**: ΔLP, radius shrink, burst quality, hallucination/omission rates.
- **String-level Effects**: Fewer loops, cleaner exits, higher prompt fidelity.
- **Safety/Neutrality**: No regression on calm prompts.

---

5) Next Steps

1. Run wobble packs through B and C configs.
2. Compare ΔLP distributions and string-level outcomes against v4b baseline.
3. If uplifts are clear, extend to live prompt sets (beyond wobble).
4. Draft continuation language if framework claim is desired.

Summary: Stage 11 doctrine is validated in patents and now baselined in practice. A/B/C experiments are the proving ground to show uplift at each doctrinal layer. Crossing these checkpoints means we can claim not just a theory, but a product-ready runtime module.