Stage■11: Revised No■Go / Go Plan — Operationally Sufficient Basin (OSB)

Context

Stage■10: Geodesic parser/executor worked in synthetic space but left phantom wells.

Stage■11 Doctrine: Warp → Detect → Denoise. Explicit warped manifold, funnel fit, phantom suppression, and denoising control stack.

Empirical proof (Latent■ARC, n=100): Stock ≈49% exact; Stage■10 Geodesic ≈64% exact; Stage■11 Denoiser = 100/100 exact, hallucination ≈0.5%, omission ≈0.2%.

Conclusion: Phantoms persist geometrically, but denoiser controls render them inert. Behavior matters more than landscape purity.

Revised OSB Checklist (Synthetic Setting)

Hard Behavior Gates (must pass)	
Accuracy exact	Target met (Latent■ARC: 100/100 exact)
Hallucination	≈ noise floor (≤0.5%)
Token drift stability	r_trend_tokens ≥ 0.90
Trajectory health	Radius ↓, SNR ↑ across descent

Soft Landscape Health (advisory)	
Phantom Index (PI)	≤ 0.30 post■warp
Margin	≥ 0.03
S_median	≥ 0.50

Stress Probes (≥3/4 should hold)	
	Half the steps
	+25% jitter
	Tap shift (–3 ↔ –2)
	20% prompt mix shift

When Tighten (LLM Integration)	
Phantom Index (PI)	≤ 0.15
Margin	≥ 0.04
S_median	≥ 0.55
Calibration	≥ 300 on■topic prompts

No■Go Criteria (Synthetic)

• Accuracy < target, or hallucination >1%.

• r_trend_tokens < 0.80 consistently (unstable drift).

• Radius increases or SNR collapses across descent.

• Any stress probe failure that also triggers a hard gate failure.

Go Criteria (Synthetic)

• Accuracy = 100/100 exact, hallucination suppressed.

• r_trend stable (≥0.90 main run, ≥0.80 under stress).

• Radius shrinks, SNR climbs.

• Stress probes pass.

Bottom Line

Clean basin not required: good■enough basin + denoiser = deterministic behavior.

Synthetic GO achieved: Stage■11 denoiser yields perfect accuracy and suppresses hallucinations to the noise floor.

Next frontier: test LLM hooks under OSB plan, applying tighter soft bars and ≥300 on■topic calibration prompts.